

4TN4 Phase 2: Image Compression Pipeline

Eric Nguyen
McMaster University
1280 Main St W, Hamilton, ON L8S 4L8
Nguyee13@McMaster.ca

Abstract

Phase 2 is an improvement on the simple lossy image compression method built in Phase 1. The image compression is still performed by using a color transformation between the RGB and YUV color spaces. However, the image decoding and inverse color transformation is performed by a convolutional neural network instead of using bilinear interpolation or bicubic interpolation methods. The reason why the input image is transformed into the YUV color space is because YUV images can be compressed at a greater factor than RGB images while preserving the same amount of digital information. This behavior is due to the fact that the red, green, and blue channels in the RGB color space are more correlated to one another compared to the luminance (Y), chrominance (U, V) channels in the YUV color space. The luminance (Y) channel also stores more meaningful information than the chrominance channels as it represents the brightness and contrast of the image which is crucial to a person's perception of an image. This fact is taken advantage of by down sampling the U and V channels twice as much as the Y channel during the encoding process. The resulting encoded YUV channel images are then fed into a convolutional neural network that is specifically trained to convert YUV channel images into a 512x512 RGB image. After testing this method on a dataset of images, the resulting peak signal-to-noise ratio and structural similarity index of this method has greatly improved compared to the phase 1 compression pipeline. The significant increase in image quality is due to the flaw found in phase 1's approach to decoding an image. Machine learning models can restore fine details and the edges of encoded images, which is where bilinear and bicubic interpolation techniques perform poorly. Although this method is much more computationally expensive than phase 1's method, the significant improvement of image quality opens the possibilities for great breakthroughs in image processing techniques.

1. Introduction

Fundamental to all forms of technology is the transmission and storage of files. There is a critical need to compress information, particularly in the modern day where data is produced at a faster rate than we can usefully process. Phase 2 is an improvement to Phase 1's implementation of an image compression pipeline by using a convolutional neural network. The flow of the pipeline remains the same, which is to compress an RGB image by first converting the image to a YUV color space. The resulting YUV image is then split into its three channels, where the chrominance (U,V) channels are down sampled twice as much as the luminance (Y) channel. The difference between Phase 1 and Phase 2's implementation lies in the image decoding or up-sampling method. Instead of using bilinear interpolation to decode the down-sampled YUV channels, a convolutional neural network is used instead. This convolutional neural network has been specifically trained and tuned to receive the down sampled YUV channel images and return a 512x512 RGB image. The reasoning behind each step will be explained with source code, mathematical equations, and diagrams.

1.1. Program Structure

The goal is to implement an image compression method without the use of external built-in functions from libraries such as OpenCV. However, NumPy is used for matrix and vector operations such as storing images and transforming color spaces. Python Image Library (PIL) was also used for easier image loading and saving. TensorFlow is also used to construct and train the convolutional neural network on the DIV2K dataset.

The structure of this image compression method is split into six steps and the modules are executed/called in the following order:

Step 1: Load the Original Image

- *Main.py*: This is the main program file that the user executes. This implements the IO pipeline for the user to input their RGB image. This file also calls the necessary functions to perform the

remaining image transformations and prints the image comparison metrics.

Step 2: Transform Image from RGB to YUV

- *Colour_conversion.py*: This module contains the *rgb_to_yuv* function which loops through each RGB pixel of the input image and extracts the corresponding YUV channel values.

Step 3: Down sample YUV Image

- *downsample.py*: Implements the down sampling of a YUV image with the *down_sample* function. Down sampling is done using the averaging of nearby pixels. This particular pipeline down samples channel Y by a factor of two and channels U, V by 4. This function returns the down-sampled Y, U, V channels.

Step 4: Pre-process Dataset for Training

- *process_data_phase1.py*: This script creates the necessary directories to access and store the processed dataset. The dataset is split into training and validation. Each image in the dataset is first resized to a uniform shape of 512x512. They are then converted to the YUV color space and down sampled to form the training/validation datasets. Each YUV channel image sets, and RGB label images are stored and saved in corresponding NumPy arrays. These NumPy arrays will be used in the CNN training process.

Step 5: Construct and Train CNN

- *train_CNN_2.py*: This script constructs and compiles the CNN model. The training data NumPy arrays generated in *process_data_phase1.py* is loaded and used to train the CNN. Once training is completed, the model is saved for future predictions in *main.py*. This module also contains function *model_predict(model, y_img, u_img, v_img)*, which receives a model and YUV channel images to return the saved model's prediction. A pyplot is also generated to display the training and validation loss during the training process.

Step 6: Load Saved Model and Predict

- *main.py*: The trained model is loaded and calls *model_predict()* from *train_CNN_2* to produce an upscaled RGB image from the down sampled YUV channel images.

Step 7: Compute PSNR

- *Psnr_mse.py*: The function *MSE_PSNR* computes the mean square error and peak signal-to-noise ratio by comparing the pixels of both the original and final decoded image. Both values of MSE and PSNR are returned when the function is called.

Step 8: Compute SSIM

- *SSIM.py*: The function *SSIM* is called from *main.py* and returns the structural similarity index

of the original RGB image and the model's RGB prediction output.

1.2. Loading the Original Image

The following code is in *main.py*, which is how the user inputs their target input RGB image file. The image is opened from the source directory using PIL and is immediately converted into a NumPy array.

```
9 #PART 0: Load original image
10 print("#PART 0: Load original image")
11 img_RGB= np.array(Image.open('picture.png'))
12 array_size = img_RGB.shape #image is currently imported as RGB
```

1.3. Color Space Transformation

Once the user inputs a valid image, *rgb_to_yuv* is immediately called from *main.py*.

```
14 #PART 1: Convert img from RGB TO YUV Color space
15 print("PART 1: Convert img from RGB TO YUV Color space")
16 img_YUV = colour_conversion.rgb_to_yuv(img_RGB)
17 #SPLIT UP CHANNELS
18 b_y = img_YUV[:, :, 0]
19 b_u = img_YUV[:, :, 1]
20 b_v = img_YUV[:, :, 2]
```

Rgb_to_yuv(rgb_img) creates an empty NumPy array with the same dimensions as the original RGB image. Note the original RGB image dimensions is Height x Width x 3 (3 color channels: R, G, B). Each pixel of the input image is looped over, and the corresponding R, G, B channel value is saved. Each corresponding YUV channel value of the pixel is calculated by multiplying specific coefficients with each RGB channel value. [1] Note that channels U, V are a linear transformation of the R, G, B channels and can potentially result in a negative value. As a result, U and V are shifted by +128 to shift the range from [-128,127] to [0,255] to match the range of an 8-bit image.

$$Y = 0.299 * R + 0.587 * G + 0.114 * B \quad (1)$$

$$U = -0.14713 * R - 0.28886 * G + 0.436 * B + 128 \quad (2)$$

$$V = 0.612 * R - 0.51499 * G - 0.10001 * B + 128 \quad (3)$$

Once each YUV value is computed, it is stored in the corresponding indices in the originally empty NumPy array *yuv_img*. Once the nested for loops are completed, *yuv_img* is returned.

```

2 import numpy as np
3
4 def rgb_to_yuv(rgb_img):
5     # Extract image dimensions
6     height, width, depth = rgb_img.shape
7     # Initialize output image as numpy arrays of zeroes
8     yuv_img = np.zeros((height, width, 3), dtype=np.float32)
9
10    for y in range(height):
11        for x in range(width):
12            # Extract RGB values from the input image
13            R, G, B = rgb_img[y, x, 0], rgb_img[y, x, 1], rgb_img[y, x, 2]
14
15            # Convert RGB to YUV
16            #***IMPORTANT: channels U,V are shifted +128 to go from [-128,127] to
17            # [0,255]; this will result in loss of color after quantization in up and downsample
18            Y = int(0.299*R + 0.587*G + 0.114*B)
19            U = int(-0.14713*R - 0.28886*G + 0.436*B+128)
20            V = int(0.615*R - 0.51499*G - 0.10001*B+128)
21
22            # Set the YUV values in the output image
23            yuv_img[y, x, 0] = Y
24            yuv_img[y, x, 1] = U
25            yuv_img[y, x, 2] = V
26
27    return yuv_img

```

1.4. Down Sampling YUV Image

Down sampling the YUV image is simply resizing each channel by a specific factor. In this case, the chrominance channels U, V are down sampled twice as much as the luminance channel Y. The reason for this design choice is because the human visual system extracts much more meaningful information from the Y channel compared to the chrominance channels. This means that the UV channels can be compressed much more without sacrificing image quality. This is the sole reason why the image was converted from RGB to YUV.

The function `downsample` is called from `main.py` with the YUV image and a `downsample_factor` of 2 as inputs. Three down sampled YUV channels are returned from this function call.

```

25 #PART 2: Downsample YUV image
26 print("PART 2: Downsample YUV image")
27 downsampled_y, downsampled_u, downsampled_v = up_down_sampling.downsample(img_YUV, downsample_factor = 2)

```

`Downsample` is a wrapper function that calls `downsample_image` to down sample all 3 YUV channels. Note how channels U and V are down sampled with a factor twice as large as Y.

```

94 def downsample(img, downsample_factor):
95     downsampled_y = downsample_image(img[:, :, 0], downsample_factor)
96     downsampled_u = downsample_image(img[:, :, 1], downsample_factor*2)
97     downsampled_v = downsample_image(img[:, :, 2], downsample_factor*2)
98     return downsampled_y, downsampled_u, downsampled_v

```

The function `downsample_image` will call two functions, `create_downsampling_kernel` and `pad_image`. `Create_downsampling_kernel` creates a numpy array with the shape of (`downsample_factor` X `downsample_factor`). `Pad_image` will pad the input image with zeroes to ensure no out of bounds issues during the convolution downsampling.

```

58 def downsample_image(img, downsample_factor):
59     input_height, input_width = img.shape[0], img.shape[1]
60
61     # Create downsampling kernel
62     kernel = create_downsampling_kernel(downsample_factor)
63
64     # Create a numpy array with downsampled dimensions filled with zeroes:
65     output_height, output_width = int(input_height/downsample_factor), int(input_width/downsample_factor)
66     downsampled_channel = np.zeros((output_height, output_width), dtype=np.uint8)
67
68     # Pad the input image to prevent out of bounds issues with kernel size mismatch with input input size
69     padded_image = pad_image(img, kernel, downsample_factor)
70

```

```

24 def create_downsampling_kernel(kernel_size):
25     kernel = np.ones((kernel_size, kernel_size), dtype=np.float32) / (kernel_size**2)
26     return kernel
27

```

A kernel is created to perform nearest neighbor averaging, which is averaging the dot product of a section of the input image and the kernel. For example, a 2x2 pixel section of the original image would be averaged to become the value of a single pixel in the down sampled image.

```

71 # Iterate over each pixel of the PADDED input image
72 for i in range(output_height):
73     for j in range(output_width):
74
75         # Extract the section of pixels that surround the currently selected pixel with the same shape of kernel
76         surrounding_pixels = padded_image[i*downsample_factor:i*downsample_factor +
77                                           kernel.shape[0], j*downsample_factor:j*downsample_factor + kernel.shape[1]]
78         # Perform element-wise multiplication between the patch and the kernel
79         convoluted_section = surrounding_pixels * kernel
80         # Sum the products of the multiplication
81         final_pixel = np.sum(convoluted_section, axis=(0, 1))
82         # Quantize final pixel to [0,255], and assign to downsampled_channel
83         final_pixel = int(final_pixel)
84         if (final_pixel < 0):
85             final_pixel = 0
86         elif (final_pixel > 255):
87             final_pixel = 255
88         downsampled_channel[i,j] = final_pixel
89
90 return downsampled_channel

```

In this code, an image is down sampled by looping through each pixel of the image. In this nested for loop, the kernel is mapped onto the current pixel and nearby pixels to compute an average value of that 2x2 section. The average of the 2x2 section is then quantized to an integer value between 0 and 255 and is copied to the output array `downsampled_channel` that represents the encoded image. On the subsequent iterations, the kernel is shifted, and a new average is computed until the entire original input image has been down sampled.




Figure 1: The kernel acts as a "sliding window" that is used to compute the average of the section of pixels.

Recall that the wrapper function `downsample()` also called `pad_image()`. This was done to ensure that the image being down sampled has dimensions divisible by the down sample factor. If the input image were to not be padded, there may be an out of bounds error where the kernel may be overlapping with pixels "outside" of the input image that do not exist.

```

23 def pad_image(img, kernel, downsample_factor):
24     # Take input dimensions and scale output dimensions
25     input_height, input_width = img.shape[0], img.shape[1]
26
27     # Compute the padding sizes; add extra padding if dimensions is not divisible by downsample factor to prevent out of bounds error
28     # If the original image height/width is not divisible by downsample factor, we must add an extra layer because we will end up
29     # convolving the kernel over pixels outside of the image that don't exist.
30
31     top_pad = kernel.shape[0] // 2
32     if (input_height % downsample_factor == 0):
33         bottom_pad = top_pad
34     else:
35         bottom_pad = top_pad + 1
36
37     left_pad = kernel.shape[1] // 2
38     if (input_width % downsample_factor == 0):
39         right_pad = left_pad
40     else:
41         right_pad = left_pad + 1
42
43     # Image is padded with zeros
44     padded_image = np.zeros((input_height + top_pad + bottom_pad, input_width + left_pad + right_pad), dtype=np.float32)
45     # Replace the non padded section of padded image with the original image, = put input image in middle of padded_image
46     padded_image[top_pad:top_pad+input_height, left_pad:left_pad+input_width] = img
47
48     return padded_image

```

Once these three down sampling steps are completed for each Y, U, V channel, the three images are returned to main.py.

```

70 return downscaled_y, downscaled_u, downscaled_v

```

1.5. Data Preprocessing

The machine learning model is only as good as the data it is trained on, which makes it crucial to preprocess the dataset. There are five main steps performed in this script:

```

299 if __name__ == '__main__':
300     #1. MAKE DIRECTORIES
301     make_dir()
302     #2. RESIZE IMAGES AND SAVE TO CORRESPONDING DIRECTORIES
303     process_data(512, 512)
304     #3. Perform colour transformations and down sampling
305     generate_train_data()
306     generate_valid_data()
307     #4 GENERATE TRAINING AND VALIDATION NUMPY ARRAYS
308     y, u, v = generate_train_np()
309     y_valid, u_valid, v_valid = generate_validation_np()
310     rgb_train = generate_RGB_train_np()
311     rgb_valid = generate_RGB_valid_np()
312     #5 SAVE DATASET NUMPY ARRAYS TO DIRECTORY
313     save_data_as_numpy_arrays(y, u, v, y_valid, u_valid, v_valid, rgb_train, rgb_valid)
314

```

1.5.1 Directory Creation

When the images are preprocessed for the model training, they will be saved to either a training or validation dataset. Each image will also be converted to the YUV color space and each down-sampled channel will be saved as a separate image.

```

32 #make directories
33 train_data_path = os.path.join(os.getcwd(), 'data', 'DIV2K_train_HR', 'resized')
34 valid_data_path = os.path.join(os.getcwd(), 'data', 'DIV2K_valid_HR', 'resized')
35
36 def make_dir():
37     try:
38         #training data
39         os.makedirs(os.path.join(train_data_path, 'y_channel', 'downsampled'), 0o666)
40         os.makedirs(os.path.join(train_data_path, 'u_channel', 'downsampled'), 0o666)
41         os.makedirs(os.path.join(train_data_path, 'v_channel', 'downsampled'), 0o666)
42         #validation data
43         os.makedirs(os.path.join(valid_data_path, 'y_channel', 'downsampled'), 0o666)
44         os.makedirs(os.path.join(valid_data_path, 'u_channel', 'downsampled'), 0o666)
45         os.makedirs(os.path.join(valid_data_path, 'v_channel', 'downsampled'), 0o666)
46         os.mkdir(os.path.join(os.getcwd(), 'data_arrays'))
47         print('1. Directories created.')
48     except Exception as e:
49         print(e)

```

1.5.2 Resize Source Images

The DIV2K dataset contains 900 images of varying dimensions which will all be resized to 512x512 images.

The reason why these images are resized is because the model being trained in this implementation requires a dataset of uniform shape.

1.5.3 Generate Training and Validation Datasets

The resized images are now transformed to the YUV color space and down-sampled. This function converts an input image to the YUV color space, and down-samples each individual channel. The color conversion and down-sampling are performed by calling functions from colour_conversion.py and downsample.py. This function is called for all images in both the resized training and validation datasets.

```

70 def generate_data_pairs(input_img, path):
71     downsample_ratio_y = 2
72     downsample_ratio_uv = downsample_ratio_y*2
73     train_data_path = path
74     img_path = os.path.join(train_data_path, input_img)
75
76     rgb_img = np.array(Image.open(img_path))
77     yuv_img = colour_conversion.rgb_to_yuv(rgb_img)
78
79     y = yuv_img[:, :, 0]
80     u = yuv_img[:, :, 1]
81     v = yuv_img[:, :, 2]
82
83     downscaled_y = downsample.downsampler(y, downsample_ratio_y)
84     downscaled_u = downsample.downsampler(u, downsample_ratio_uv)
85     downscaled_v = downsample.downsampler(v, downsample_ratio_uv)
86
87
88     return downscaled_y, downscaled_u, downscaled_v

```

1.5.4 Convert Training and Validation Images to NumPy

TensorFlow models cannot be trained on images directly. These images must be converted to NumPy arrays to train the model. The generate_train_np() function iterates through each down-sampled Y, U, V image set in the training and validation data directories. This particular section of code displays the initialization of the NumPy array 'y_down_train_np'. Each Y channel image of the training dataset will be converted to a NumPy array representation and is appended to this array. This is repeated for all Y, U, V color channels and the resized RGB images in both training and validation datasets.

```

def generate_train_np():
    y_down_train = (os.path.join(os.getcwd(), 'data', 'DIV2K_train_HR', 'resized', 'y_channel', 'downsampled'))
    #create a list containing all filenames
    image_filenames = [filename for filename in os.listdir(y_down_train) if filename.endswith('.png')]

    down_y_img = Image.open(os.path.join(y_down_train, image_filenames[0]))
    d_width, d_height = down_y_img.size
    #initialize empty np array
    y_down_train_np = np.empty((len(image_filenames), d_height, d_width))

    for i, image_filename in enumerate(image_filenames):
        img = Image.open(os.path.join(y_down_train, image_filename))
        y_down_train_np[i] = np.array(img)

```

1.5.5 Save NumPy Arrays to Directory

Once the entire dataset has been converted and saved to their respective NumPy array, these arrays will be saved as a .npz file in the data_arrays directory.

```
277 def save_data_as_numpy_arrays(y,u,v,y_valid,u_valid,v_valid,rgb_train,rgb_valid):
278     y_train = y
279     u_train = u
280     v_train = v
281
282     np.save(os.path.join(os.getcwd(), 'data_arrays', 'y_train'), y_train)
283     np.save(os.path.join(os.getcwd(), 'data_arrays', 'u_train'), u_train)
284     np.save(os.path.join(os.getcwd(), 'data_arrays', 'v_train'), v_train)
285
286     np.save(os.path.join(os.getcwd(), 'data_arrays', 'y_valid'), y_valid)
287     np.save(os.path.join(os.getcwd(), 'data_arrays', 'u_valid'), u_valid)
288     np.save(os.path.join(os.getcwd(), 'data_arrays', 'v_valid'), v_valid)
289     np.save(os.path.join(os.getcwd(), 'data_arrays', 'rgb_train'), rgb_train)
290     np.save(os.path.join(os.getcwd(), 'data_arrays', 'rgb_valid'), rgb_valid)
291
292     print('9. All numpy arrays saved to main working directory.')
```

1.6. Construct and Train CNN

A convolutional neural network is a type of deep neural network commonly used for image classification and recognition. A convolutional neural network consists of a series of layers that perform different operations on the input data. The purpose of these connected layers is for the neural network to learn increasingly complex representations of the input data. Ultimately, the goal of a neural network is to eventually learn the relationships between the input training data and the ground truths to make accurate predictions on new data. In this case, the CNN is being trained to determine the relationship between the down-sampled Y, U, V channel images and the original RGB image.




Figure 2 A simple 1D CNN model

1.6.1 Data Loading

Now that the original dataset has been preprocessed and saved as NumPy arrays, it is now possible to train a CNN. First, the data arrays are loaded and divided by 255.0 to normalize the dataset. Normalization is crucial for reducing complexity which ultimately increases the possibility of the model to converge. The following code is found in train_CNN.py.

```
43 y_train = np.load(os.path.join(os.getcwd(), 'data_arrays', 'y_train.npy'))/255.0
44 u_train = np.load(os.path.join(os.getcwd(), 'data_arrays', 'u_train.npy'))/255.0
45 v_train = np.load(os.path.join(os.getcwd(), 'data_arrays', 'v_train.npy'))/255.0
46 rgb_train = np.load(os.path.join(os.getcwd(), 'data_arrays', 'rgb_train.npy'))/255.0
47
48 rgb_train = rgb_train[:, :, :]
49
50 y_valid = np.load(os.path.join(os.getcwd(), 'data_arrays', 'y_valid.npy'))/255.0
51 u_valid = np.load(os.path.join(os.getcwd(), 'data_arrays', 'u_valid.npy'))/255.0
52 v_valid = np.load(os.path.join(os.getcwd(), 'data_arrays', 'v_valid.npy'))/255.0
53 rgb_valid = np.load(os.path.join(os.getcwd(), 'data_arrays', 'rgb_valid.npy'))/255.0
```

1.6.2 CNN Model Design

The convolutional neural network is constructed and compiled in the code below:

```
72 def construct_model():
73     # Define the input shape for the model
74     y_input_shape = (256, 256, 1)
75     uv_input_shape = (128, 128, 1)
76     # Define the input layers for the model
77     y_input = Input(shape=y_input_shape, name='y_input')
78     u_input = Input(shape=uv_input_shape, name='u_input')
79     v_input = Input(shape=uv_input_shape, name='v_input')
80     # First block of the model: convolve Y channel with 32 filters of size 3x3
81     y_conv = Conv2D(32, (3, 3), padding='same')(y_input)
82     y_conv = BatchNormalization()(y_conv)
83     y_conv = Activation('relu')(y_conv)
84     # Second block of the model: concatenate U and V channels and convolve with 32 filters of size 3x3
85     uv_concat = Concatenate()([u_input, v_input])
86     uv_conv = Conv2D(32, (3, 3), padding='same')(uv_concat)
87     uv_conv = BatchNormalization()(uv_conv)
88     uv_conv = Activation('relu')(uv_conv)
89     # Third block of the model: upsample U and V channels and concatenate with Y channel
90     uv_upsample = UpSampling2D(size=(2, 2))(uv_conv)
91     merged = Concatenate()([y_conv, uv_upsample])
92     # Fourth block of the model: convolve merged channels with 64 filters of size 3x3
93     conv = Conv2D(64, (3, 3), padding='same')(merged)
94     conv = BatchNormalization()(conv)
95     conv = Activation('relu')(conv)
96     conv = Dropout(0.3)(conv)
97     # Fifth block of the model: upsample and convolve with 64 filters of size 3x3
98     upsample = UpSampling2D(size=(2, 2))(conv)
99     conv = Conv2D(64, (3, 3), padding='same')(upsample)
100    conv = BatchNormalization()(conv)
101    conv = Activation('relu')(conv)
102    conv = Dropout(0.3)(conv)
103    # Sixth block of the model: apply a final convolutional layer with 3 filters of size 3x3
104    output = Conv2D(3, (3, 3), padding='same', activation='linear')(conv)
105    # Define the model with the input and output layers
106    model = Model(inputs=[y_input, u_input, v_input], outputs=output)
107    # Compile the model
108    optimizer = keras.optimizers.RMSprop(lr=0.001, rho=0.9, epsilon=None, decay=0.0)
109    model.compile(optimizer=optimizer, loss='MAE')
110
111    return model
```

Model: "model"			
Layer (type)	Output Shape	Param #	Connected to
u_input (InputLayer)	[(None, 128, 128, 1, 0)]		[]
v_input (InputLayer)	[(None, 128, 128, 1, 0)]		[]
concatenate (Concatenate)	(None, 128, 128, 2) 0		['u_input[0][0]', 'v_input[0][0]']
y_input (InputLayer)	[(None, 256, 256, 1, 0)]		[]
conv2d_1 (Conv2D)	(None, 128, 128, 12, 2432)		['concatenate[0][0]']
conv2d (Conv2D)	(None, 256, 256, 12, 1280)		['y_input[0][0]']
batch_normalization_1 (Batch Normalization)	(None, 128, 128, 12, 512)		['conv2d_1[0][0]']
conv2d_2 (Conv2D)	(None, 256, 256, 64, 147520)		['concatenate_1[0][0]']
batch_normalization_2 (Batch Normalization)	(None, 256, 256, 64, 256)		['conv2d_2[0][0]']
activation_2 (Activation)	(None, 256, 256, 64, 0)		['batch_normalization_2[0][0]']
dropout (Dropout)	(None, 256, 256, 64, 0)		['activation_2[0][0]']
up_sampling2d_1 (UpSampling2D)	(None, 512, 512, 64, 0)		['dropout[0][0]']
conv2d_4 (Conv2D)	(None, 512, 512, 3) 1731		['up_sampling2d_1[0][0]']
Total params: 154,243			
Trainable params: 153,603			
Non-trainable params: 640			

Figure 3 Model Summary displaying each layer.

This particular model requires three inputs: a Y channel array of shape (256,256,1), U channel and V channel arrays of shape (128,128,1). These input specifications are created on lines 77-79.

The remaining section of code consists of the model architecture, model optimizer, and model compilation.

1.6.2.1 Model Layer Breakdown

Each layer of the neural network has a convolution operation performed. Convolution layers apply a set of filters to the input data to extract features. This is very similar to the use of kernels in the down-sampling process in section 1.4. The stacking of convolutional layers allows the neural network to learn increasingly complex features by combining the simple features learned in previous layers. Each convolution layer contains weights which are used to recognize patterns in the data. During training, the model will take the input data and produce predictions based on its current set of weights. The error between the prediction and true output is calculated using a loss function. These weights will be adjusted to minimize the loss function. If the model converges, it will be able to map out the relationship between the YUV channel image size and color values to an RGB image size and color values.

1.6.2.1.1 Layer 1: Y channel Convolution

```
80 # First block of the model: convolve Y channel with 32 filters of size 3x3
81 y_conv = Conv2D(128, (3, 3), padding='same')(y_input)
82 y_conv = BatchNormalization()(y_conv)
83 y_conv = Activation('relu')(y_conv)
```

This first layer handles the Y channel input. It initially applies a 2D convolution the `y_input` with 128 filters and a kernel size of (3,3), and will return an output of the same shape of `y_input`.

It will then apply a batch normalization on the convolution layer's output to reduce covariate shift. It is essentially re-centering and re-scaling the output by maintaining the batch's mean output close to 0 and output standard deviation close to 1. This will allow faster and more stable training.

The rectified linear unit activation function 'ReLU' is applied to the batch normalization layer's output. ReLU converts all negative values to 0. The purpose of this activation function is to allow neural networks to learn and generalize better. This was used because it is computationally efficient and introduced nonlinearity to the model. It is crucial to introduce nonlinearity as the model would simply be a linear regression model if ReLU was not used. Nonlinear models perform better at recognizing any complex nonlinear patterns that may exist in the images.

1.6.2.2 Layer 2: U, V channel concatenation and convolution

```
84 # Second block of the model: concatenate U and V channels and convolve with 32 filters of size 3x3
85 uv_concat = Concatenate()([u_input, v_input])
86 uv_conv = Conv2D(128, (3, 3), padding='same')(uv_concat)
87 uv_conv = BatchNormalization()(uv_conv)
88 uv_conv = Activation('relu')(uv_conv)
```

The U and V input layers are concatenated and the exact same operations performed in Layer 1 are performed on the U, V input layers.

1.6.2.3 Layer 3: UV Up-sampling and Y channel Concatenation

```
89 # Third block of the model: upsample U and V channels and concatenate with Y channel
90 uv_upsample = UpSampling2D(size=(2,2))(uv_conv)
91 merged = Concatenate()([y_conv, uv_upsample])
```

This layer up-samples the U and V channels to match the shape of the Y channel. The UV layer and Y input layer are then concatenated once they are the same size. This concatenation is crucial as convolutional layers operate on the image data as a whole, and not just on individual pixels. This will allow the neural network to learn the color relationships between the YUV and RGB color spacing which is crucial for the inverse color transformation operation.

1.6.2.4 Layer 4: Convolution

```
92 # Fourth block of the model: convolve merged channels with 64 filters of size 3x3
93 conv = Conv2D(64, (3, 3), padding='same')(merged)
94 conv = BatchNormalization()(conv)
95 conv = Activation('relu')(conv)
96 conv = Dropout(0.3)(conv)
```

The concatenated Y, U, V layers are convolved with 64 filters and a kernel size of (3,3). Dropout is a regularization technique that randomly sets values in the output to zero during each training epoch. Dropout was used to prevent the model from overfitting to the training dataset. This ultimately forces the model to learn more robust and generalizable features instead of neurons becoming too specialized to particular features. In this case, a dropout with a rate of 0.3 sets 30% of the values to zero.

1.6.2.5 Layer 5: Up-sample and Convolution

```
97 # Fifth block of the model: upsample and convolve with 64 filters of size 3x3
98 upsample = UpSampling2D(size=(2,2))(conv)
99 conv = Conv2D(64, (3, 3), padding='same')(upsample)
100 conv = BatchNormalization()(conv)
101 conv = Activation('relu')(conv)
102 conv = Dropout(0.3)(conv)
```

The concatenated Y, U, V layers are now up sampled to match the desired image size of 512x512. Convolution, batch normalization, ReLU activation, and dropout are also performed.

1.6.2.6 Layer 6: Final Convolution

```
103 # Sixth block of the model: apply a final convolutional layer with 3 filters of size 3x3
104 output = Conv2D(3, (3, 3), padding='same', activation='linear')(conv)
```

The final convolution layer is performed with an activation function named 'linear'. This function does not change the input values and allows the model to output any real number.

$$f(x) = x \quad (4)$$

There are other activation functions such as sigmoid or previously mentioned ReLU. For example, the sigmoid function produces an output with a probability value between 0 and 1. This is particularly used in binary classification applications.

$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)} \quad (4)$$

1.6.2.7 Optimizer and Loss Function

The model is now created with the Keras API. An optimizer and the loss function are also specified.

```
107 model = Model(inputs=[y_input, u_input, v_input], outputs=output)
```

The loss function measures how accurate the model's predictions are given the input data.

An optimizer in machine learning is an algorithm that adjusts the weights and biases to minimize the loss function during training. There are many different types of optimizers, but they all control how fast the model will converge, the model's ability to generalize to new data, and each performs differently to outliers in training data.

In this case, a Ro RMSprop (Root Mean Square) optimizer with a learning rate of 0.001, a decay rate of 0.0, momentum parameter of 0.9, and epsilon constant of None. The RMSprop optimizer is a variant of the stochastic gradient descent algorithm. This optimizer uses a moving average of the squared gradients to normalize the gradient. The goal of this optimizer is to prevent oscillations in the optimization process and to speed up convergence time.

The optimizer updates the weight in this fashion:

$$\text{weight}_{\text{new}} = \text{weight}_{\text{old}} - \frac{\text{learning rate}}{\sqrt{\text{Root Mean Square}(\text{gradient})}} * \text{gradient} \quad (5)$$

```
109 optimizer = keras.optimizers.RMSprop(lr=0.001, rho=0.9, epsilon=None, decay=0.0)
```

The loss function used in this model is the Mean Square Error (MSE) loss function. MSE is a very common loss function used in regression problems. In this case, the MSE loss function would calculate the squared difference between the predicted RGB image and the true RGB input image. The loss function measures the average of the squares of errors meaning that it gives more weight to larger errors.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (I_i - \hat{I}_i)^2 \quad (6)$$

```
110 model.compile(optimizer=optimizer, loss='MSE')
```

1.6.3 Model Training

Now that the model has now been constructed and compiled, it can now be trained with the training and validation data. The model is fed with the Y, U, V and

RGB channel images as training and validation data in model.fit(). The batch size represents the number of training images that are passed through each pass of the model. Batch size is typically chosen as a power of two as most hardware components are optimized for power of two sizes. Epochs represents the number of times the entire dataset will be processed.

There is an early_stopping_callback which allows the model to stop training if the model's validation loss does not improve after 10 consecutive epochs. This was set to reduce overfitting.

The trained model and it's history are now saved to the main directory.

```
159 #allow early stopping of model.fit to prevent overfitting
160 early_stopping_callback = EarlyStopping(monitor='val_loss', patience=10)
161 history = model.fit([y_train, u_train, v_train], rgb_train,
162                   batch_size=4,
163                   epochs=100,
164                   callbacks=[early_stopping_callback, tf.keras.callbacks.History()],
165                   validation_data = ([y_valid, u_valid, v_valid], rgb_valid))
166 np.save('model_history.npy', history.history)
167 model.save('trained_model_rgb_out_v7')
```

1.7. Load Saved Model and Predict

Now that the model has been trained, it can now be loaded and fed new input data to form a prediction. The down-sampled YUV channels from the test image and the saved model are now fed into model_predict(). The output prediction is a NumPy array and is displayed using PIL.Image.fromarray and .show(). The predicted image is also saved to the main directory.

```
58 prediction_np = train_CNN.model_predict(loaded_model, y, u, v)
59 prediction = Image.fromarray(prediction_np)
60 prediction.show()
61
62 prediction.save(os.path.join(os.getcwd(), 'final_image_phase2.png'))
```

1.8. Calculate Peak Signal-to-Noise Ratio

Peak signal-to-noise ratio represents the measurement of quality between the original and compressed image. The PSNR is greater as quality of the reconstructed image increases. Mean squared prediction error represents the cumulative error between the compressed and original error. PSNR represents the peak error. MSE and PSNR can be calculated using the following formulae:

$$\text{MSE} = \frac{1}{\text{Width} * \text{Height} * 3} \sum_{i,j,c} (I_{i,j,c} - \hat{I}_{i,j,c})^2 \quad (7)$$

$$\text{PSNR} = 10 * \log_{10} \frac{255^2}{\text{MSE}} \quad (8)$$


```

6 def MSE_PSNR(imageA, imageB):
7     MSE = np.sum((imageA - imageB) ** 2)
8     MSE /= (imageA.shape[0] * imageA.shape[1])
9     PSNR = 10*math.log10(255**2/MSE)
10    print("PSNR = ",PSNR)
11    return MSE, PSNR

```

1.9. Calculate Structural Similarity Index

The Structural Similarity Index (SSIM) is a metric used to compare two images that mimic the human visual system. It is different from PSNR as it places greater emphasis on structural properties in images such as edges and textures. It also places less emphasis on changes in color and brightness.

$$SSIM(x, y) = \frac{(2\mu_x\mu_y + c_1)(2\sigma_{xy} + c_2)}{(\mu_x^2 + \mu_y^2 + c_1)(\sigma_x^2 + \sigma_y^2 + c_2)} \quad (9)$$

where:

μ_x is the mean of x

μ_y is the mean of y

σ_x^2 is the variance of x

σ_y^2 is the variance of y

σ_{xy} is the covariance of x and y

$c_1 = (k_1L)^2$

$c_2 = (k_2L)^2$

The parameters k_1 and k_2 are constants that adjust the relative impact of the luminance and contrast terms in equation 9. K_1 and K_2 are set as 0.01 and 0.03 by default. (Nilsson & Akenine-Möller, 2020)

```

1 import numpy as np
2 def SSIM(image_A, image_B):
3     #K_1, K_2 from https://arxiv.org/pdf/2006.13846.pdf
4     K_1 = 0.01
5     K_2 = 0.03
6     L = 255 #8 bit image 0<->255
7     mean_A = np.mean(image_A)
8     mean_B = np.mean(image_B)
9     var_A = np.var(image_A)
10    var_B = np.var(image_B)
11    #need to flatten() to turn each image into a feature vector
12    #take element at [0][1] because we want covar AB,
13    #covar_AB is a 2x2 matrix:
14    # var(A), cov(A,B)
15    # cov(B,A), var(B)
16    covar_AB = np.cov(image_A.flatten(), image_B.flatten())[0][1]
17    C_1 = (K_1*L)**2
18    C_2 = (K_2*L)**2
19    numerator = (2*mean_A*mean_B + C_1)*(2*covar_AB + C_2)
20    denominator = (mean_A**2+mean_B**2+C_1)*(var_A+var_B+C_2)
21
22    SSIM = numerator/denominator
23    return SSIM
24

```

Please note that the SSIM input images must be converted to grayscale.

```

73 test_img_grayscale = rgb_to_grayscale.RGB_to_GRAYSCALE(img_RGB[:, :, :3])
74 prediction_grayscale = rgb_to_grayscale.RGB_to_GRAYSCALE(prediction_np)
75
76 SSIM_score = SSIM.SSIM(test_img_grayscale, prediction_grayscale)
77 print('SSIM = {}'.format(SSIM_score))

```

The predicted and test images are converted to grayscale using this function:

```

1 import numpy as np
2 def RGB_to_GRAYSCALE(image):
3     height, width, channels = image.shape
4     # Initialize empty np array
5     gray_image = np.zeros((height, width), dtype=np.uint8)
6
7     # Iterate over each pixel in the image and compute its grayscale value
8     for i in range(height):
9         for j in range(width):
10            # Get the red, green, and blue values of the pixel
11            r, g, b = image[i, j]
12            # Compute the grayscale value using the luminance method
13            y = 0.2126*r + 0.7152*g + 0.0722*b
14            # Store the grayscale value in the new image array
15            gray_image[i, j] = int(y)
16    return gray_image

```

2. Experimental Results

2.1. Training and Validation Loss

The initial CNN model constructed contained fewer convolution layers with less filters, did not use the Dropout technique, and the Adam optimizer. Its performance was comparably worse during training and when testing on sample images.




Figure 4 Initial CNN Model Training Performance

The training loss, PSNR, and SSIM metrics all improved after including more convolution layers, using the Dropout technique, and the RMSprop optimizer in the CNN model. It is clear that the RMSprop optimizer greatly reduced the oscillation in the validation loss.




Figure 5 Final CNN Model Training Performance

2.2. Image Quality Comparison

This image compression pipeline was tested on 5 images. The images are fairly accurate visually; however, the loss of digital information is clear when the image is zoomed in. This loss of information can be improved by using another down sampling algorithm. The loss of information is reflected in the PSNR values. However, the SSIM values are all very close to the maximum value of 1 implying that the semantic information is retained.

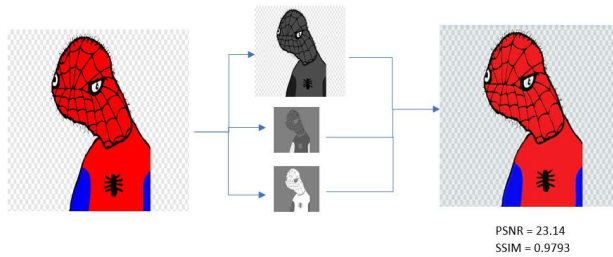


Figure 6 Test Image 1




Figure 7 Test Image 2




Figure 8 Test Image 3




Figure 9 Test Image 4




Figure 10 Test Image 5

2.3. File Size Comparison

The compressed file size of the test image dress_512x512 is 518 KB, and the original file size is 1028 KB. The resulting compression rate is approximately 50.38%.

$$\text{Compression Ratio} = \frac{(\text{size before compression})}{\text{size after compression}} \quad (9)$$

$$= \frac{518}{1028}$$

$$\cong 50.38\%$$

3. Conclusion

The goal of this image compression pipeline has demonstrated that images can be significantly compressed while preserving most of the semantic information. This pipeline demonstrated that the use of different color spaces such as the YUV color space can greatly improve image compression rates. The implementation of the convolutional neural network also made a significant increase in image quality especially around the edges of the image compared to the bilinear interpolation method used in phase 1. This use of the convolutional neural network demonstrated how machine learning will be used to make significant advances in image processing.

References

- [1] ITU-R. (2011, March 8). Recommendation ITU-R BT.2020-2. International Telecommunication Union. ,from https://www.itu.int/dms_pubrec/itu-r/rec/bt/R-REC-BT.601-7-201103-I!!PDF-E.pdf
- [2] Li, Z.-N., Drew, M. S., & Liu, J. (2021). *Fundamentals of Multimedia* (3rd ed.). Springer.
- [3] Nilsson, J., & Akenine-Möller, T. (2020, June 29). *Understanding SSIM*. arXiv.org. Retrieved April 15, 2023, from <https://arxiv.org/abs/2006.13846>